

实验三报告

关卡一：openGauss 数据库的编译和安装

1. 关卡验证

步骤 1 首先需要对数据库状态进行验证。

```
[omm@opengauss01 openGauss-server]$ gs_ctl status
```

(截图语句和执行结果)

```
[omm@opengauss01 ~]$ gs_ctl status
[2022-12-06 14:50:31.096][225310][][gs_ctl]: gs_ctl status,datadir is /opt/software/openGauss/data
gs_ctl: server is running (PID: 225266)
/opt/software/openGauss/bin/gaussdb "-D" "/opt/software/openGauss/data"
[omm@opengauss01 ~]$
```

步骤 2 对数据库进程进行截图验证，需包含数据库服务器的主机名。

```
[omm@opengauss01 openGauss-server]$ ps -ef|grep omm
```

(截图语句和执行结果)

```
[omm@opengauss01 ~]$ ps -ef|grep omm
root      225072    5175  0 14:47 pts/0    00:00:00 su - omm
omm       225073    225072  0 14:47 pts/0    00:00:00 -bash
omm       225266      1  1 14:48 pts/0    00:00:01 /opt/software/openGauss/bin/gaussdb -D /opt/software/openGauss/data
omm       225319    225073  0 14:51 pts/0    00:00:00 ps -ef
omm       225320    225073  0 14:51 pts/0    00:00:00 grep --color=auto omm
[omm@opengauss01 ~]$
```

关卡二：openGauss 数据导入及基本操作

1. 关卡验证

步骤 12 登录数据库验证

```
[omm@opengauss01 dbgen]$ gsql -d tpch -p 5432 -r
tpch=# select count(*) from supplier;
```

(截图语句和执行结果)

```
[omm@opengauss01 dbgen]$ gsql -d tpch -p 5432 -r
gsql ((GaussDB Kernel V500R002C00 build b2ff10be) compiled at 2022-12-06 14:40:30 commit 0 last mr debug)
Non-SSL connection (SSL connection is recommended when requiring high-security)
Type "help" for help.

tpch=# select count(*) from supplier;
count
-----
10000
(1 row)

tpch=#
```

步骤 21 登录数据库进行验证

```
[omm@opengauss01 ~]$ gsql -d tpch -p 5432 -r
tpch=# \dt
```

(截图语句和执行结果)

```
tpch=# \dt
```

Schema	Name	Type	Owner	Storage
public	address_dimension	table	omm	{orientation=row,compression=no}
public	customer	table	omm	{orientation=row,compression=no}
public	date_dimension	table	omm	{orientation=row,compression=no}
public	lineitem	table	omm	{orientation=row,compression=no}
public	litemail_orders	table	omm	{orientation=row,compression=no}
public	nation	table	omm	{orientation=row,compression=no}
public	orders	table	omm	{orientation=row,compression=no}
public	part	table	omm	{orientation=row,compression=no}
public	partsupp	table	omm	{orientation=row,compression=no}
public	region	table	omm	{orientation=row,compression=no}
public	supplier	table	omm	{orientation=row,compression=no}
public	user_dimension	table	omm	{orientation=row,compression=no}

(12 rows)

步骤 22 查询 customer 表的数据

```
tpch=# select * from customer limit 10;
```

(截图语句和执行结果)

```
tpch=# select * from customer limit 10;
```

c_custkey	c_name	c_comment	c_address	c_nationkey	c_phone	c_acctbal	c_mktsegment
1	Customer#000000001	IVhziApeRb ot,c,E		15	25-989-741-2988	711.56	BUILDING
2	Customer#000000002	XStf4,NCwDVaWNeStEgwmfRchLXak		13	23-768-687-3665	121.65	AUTOMOBILE
3	Customer#000000003	hely ironic theodolites integrate boldly: caref		1	11-719-748-3364	7498.12	AUTOMOBILE
4	Customer#000000004	lyly ironic, even instructions. express foxes detect stlyly. blithely even accounts abov		4	14-128-198-5944	2866.83	MACHINERY
5	Customer#000000005	, regular ideas sleep final accou		3	13-758-942-6364	794.47	HOUSEHOLD
6	Customer#000000006	have to unwind. foxes cajole accor		20	38-114-968-4951	7638.57	AUTOMOBILE
7	Customer#000000007	sits boost according to the stlyly bold packages. final accounts cajole requests. furious		18	28-198-982-9759	9561.95	AUTOMOBILE
8	Customer#000000008	, express theodolites. express, even pinto beans among the exp		17	27-147-574-9335	8819.74	BUILDING
9	Customer#000000009	regular theodolites kindle blithely courts. carefully even theodolites haggle stlyly along the ide		8	18-338-986-3675	8324.87	FURNITURE
10	Customer#000000010	cording to the requests wake thinly excuses: pending requests haggle furiousl		5	15-741-346-9878	2753.54	HOUSEHOLD

(10 rows)

2. 思考题

数据初始化中出现了 TPC-H，这是什么？

TPC-H 是事务处理性能委员会制定的基准程序之一，是针对随机查询/商业智能处理能力的决策支持基准，其主要目的是评价特定查询的决策支持能力，强调服务器在数据挖掘、分析处理方面的能力。该基准由一套面向业务的即时查询（ad-hoc）和并发数据修改组成，模拟了决策支持系统中的数据库操作，测试数据库系统复杂查询的响应时间，以每小时执行的查询数作为度量指标。

关卡三：openGauss 的 AI4DB 特性应用

1. 关卡验证

(1) 使用 X-Tuner 进行参数优化

步骤 2 在原来 CloudShell 连接窗口中查看 queries01.log。

```
[omm@opengauss01 ~]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries01.log
```

（截图执行语句和结果）

```
[omm@opengauss01 ~]$ gsql -d tpch -p 5432 -r -f /opt/software/tpch-kit/dbgen/queries/queries.sql > /opt/software/tpch-kit/dbgen/queries/queries01.log
^[[3~[omm@opengauss01 ~]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries01.log
15      |      888 | 6757715.99
17      |      861 | 6468575.72
18      |      964 | 726687.48
25      |      892 | 6781457.95
29      |      948 | 7158885.65
30      |      989 | 6888435.13
31      |      922 | 6886578.18
(7 rows)

total time: 1287247 ms
[omm@opengauss01 ~]$
```

步骤 3 切换至 root 用户，执行 X-Tuner 进行参数建议优化

```
[omm@opengauss01 ~]$ exit
[root@opengauss01 xtuner]# gs_xtuner recommend --db-name tpch --db-user omm --port 5432
--host 127.0.0.1 --host-user omm
```

（截图执行语句和结果）

name	recommend	min	max	restart
default_statistics_target	1000	100	1000	False
effective_cache_size	21602334	186756	21602334	False
effective_io_concurrency	200	150	250	False
enable_mergejoin	off	0	1	False
enable_nestloop	off	0	1	False
max_connections	370	50	741	True
max_prepared_transactions	370	50	741	True
max_process_memory	28803112	22402420	28803112	True
random_page_cost	1.0	1.0	2.0	False
shared_buffers	186756	186760	214772	True
wal_buffers	5836	2048	5836	True

[root@opengauss01 xtuner]#

步骤 6 获取参数值

```
[omm@opengauss01 ~]$ cd /opt/software/openGauss/data
[omm@opengauss01 data]$ cat postgresql.conf|grep -E
'shared_buffers|max_connections|effective_cache_size|effective_io_concurrency|wal_buffers|rando
m_page_cost|default_statistics_target'
```

（截图执行语句和结果）

```
[omm@opengauss01 data]$ cat postgresql.conf|grep -E 'shared_buffers|max_connections|effective_cache_size|effective_io_concurrency|wal_buffers|random
_page_cost|default_statistics_target'
max_connections = 370                                # (change requires restart)
# Note: Increasing max_connections costs ~400 bytes of shared memory per
shared_buffers = 187388                                # min 128kB
bulk_write_ring_size = 268                            # for bulkload, max shared_buffers
#standby_shared_buffers_fraction = 0.3                #control shared buffers use in standby, 0.1-1.0
effective_io_concurrency = 200                        # 1-1000; 0 disables prefetching
wal_buffers = 5855                                    # min 32kB
random_page_cost = 1                                # same scale as above
effective_cache_size = 21602940
default_statistics_target = 1000                      # range 1-10000
# max_locks_per_transaction * (max_connections + max_prepared_transactions)
```

步骤 7 再次执行步骤 2，对比优化前的执行时间。

（截图执行语句和结果）

```
[omm@opengauss01 data]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries01.log
13      |      888 | 6737713.99
17      |      861 | 6460573.72
18      |      964 | 7236687.40
23      |      892 | 6701457.95
29      |      948 | 7158866.63
30      |      909 | 6808436.13
31      |      922 | 6806670.18
(7 rows)

total time: 1174541 ms
[omm@opengauss01 data]$
```

优化前的时间为 1207247ms，优化后的时间为 1174541ms，比优化前少了 32706ms。

步骤 8 【附加题】有兴趣的同学可以尝试并截图记录于此。

（截图执行语句和结果）

```
[root@opengauss01 xtuner]# su - omm
```

```
[omm@opengauss01 ~]$ gs_guc set -D /opt/software/openGauss/data/ -c
"default_statistics_target = 1000" -c "effective_cache_size = 21602334" -c
"effective_io_concurrency = 200" -c "max_connections = 370" -c "max_prepared_transactions =
370" -c "max_process_memory=28803112" -c "random_page_cost = 1" -c "shared_buffers =
186756" -c "wal_buffers = 5836" -c "enable_mergejoin=off" -c "enable_nestloop=off"
```

```
[omm@opengauss01 data]$ cat postgresql.conf|grep -E 'shared_buffers|max_connections|effective_cache_size|effective_io_concurrency|wal_buffers|random_
page_cost|default_statistics_target'
max_connections = 370 # (change requires restart)
# Note: Increasing max_connections costs ~400 bytes of shared memory per
shared_buffers = 186756 # min 128kB
bulk_write_ring_size = 2048 # for bulkload, max shared_buffers
#standby_shared_buffers_fraction = 0.3 #control shared buffers use in standby, 0.1-1.0
effective_io_concurrency = 200 # 1-1000; 0 disables prefetching
wal_buffers = 5836 # min 32kB
random_page_cost = 1 # same scale as above
effective_cache_size = 21602334
default_statistics_target = 1000 # range 1-10000
# max_locks_per_transaction * (max_connections + max_prepared_transactions)
```

```
[omm@opengauss01 data]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries01.log
13 | | 888 | 6737713.99
17 | | 861 | 6460573.72
18 | | 964 | 7236687.40
23 | | 892 | 6701457.95
29 | | 948 | 7158866.63
30 | | 909 | 6808436.13
31 | | 922 | 6806670.18
(7 rows)

total time: 321901 ms
```

优化前的时间为 1207247ms，优化后的时间为 321901ms，比优化前少了 885346ms。

(2) Index-advisor: 索引推荐

步骤 4 使用 explain，对该 SQL 加以分析

```
tpch=# EXPLAIN
SELECT ad.province AS province, SUM(o.actual_price) AS GMV
FROM litemall_orders o,
     address_dimension ad,
     date_dimension dd
WHERE o.address_key = ad.address_key
     AND o.add_date = dd.date_key
     AND dd.year = 2020
     AND dd.month = 3
GROUP BY ad.province
ORDER BY SUM(o.actual_price) DESC;
```

(截图执行语句和结果)


```

tpch=# EXPLAIN
tpch=# SELECT ad.province AS province, SUM(o.actual_price) AS GMV
tpch=# FROM litemall_orders o,
tpch=# address_dimension ad,
tpch=# date_dimension dd
tpch=# WHERE o.address_key = ad.address_key
tpch=# AND o.add_date = dd.date_key
tpch=# AND dd.year = 2020
tpch=# AND dd.month = 3
tpch=# GROUP BY ad.province
tpch=# ORDER BY SUM(o.actual_price) DESC;

-----
QUERY PLAN
-----
Sort (cost=4593.80..4593.88 rows=31 width=47)
  Sort Key: (sum(o.actual_price)) DESC
  -> HashAggregate (cost=4592.72..4593.03 rows=31 width=47)
    Group By Key: ad.province
    -> Hash Join (cost=4354.43..4585.97 rows=1351 width=15)
      Hash Cond: (ad.address_key = o.address_key)
      -> Seq Scan on address_dimension ad (cost=0.00..188.02 rows=8002 width=14)
      -> Hash (cost=4337.54..4337.54 rows=1351 width=9)
        -> Hash Join (cost=1031.78..4337.54 rows=1351 width=9)
          Hash Cond: (o.add_date = dd.date_key)
          -> Seq Scan on litemall_orders o (cost=0.00..3041.00 rows=100000 width=13)
          -> Hash (cost=1031.76..1031.76 rows=2 width=4)
            -> Seq Scan on date_dimension dd (cost=0.00..1031.76 rows=2 width=4)
              Filter: ((year = 2020) AND ((month)::bigint = 3))

(14 rows)

tpch=#

```

步骤 10 使用 explain, 对该 SQL 加以分析

```

tpch=# EXPLAIN
SELECT ad.province AS province, SUM(o.actual_price) AS GMV
FROM litemall_orders o,
address_dimension ad,
date_dimension dd
WHERE o.address_key = ad.address_key
AND o.add_date = dd.date_key
AND dd.year = 2020
AND dd.month = 3
GROUP BY ad.province
ORDER BY SUM(o.actual_price) DESC;

```

(截图执行语句和结果)

```

tpch=# set enable_hypo_index = on;
SET
tpch=# EXPLAIN
tpch=# SELECT ad.province AS province, SUM(o.actual_price) AS GMV
tpch=# FROM litemall_orders o,
tpch=# address_dimension ad,
tpch=# date_dimension dd
tpch=# WHERE o.address_key = ad.address_key
tpch=# AND o.add_date = dd.date_key
tpch=# AND dd.year = 2020
tpch=# AND dd.month = 3
tpch=# GROUP BY ad.province
tpch=# ORDER BY SUM(o.actual_price) DESC;

-----
QUERY PLAN
-----
Sort (cost=3579.58..3579.65 rows=31 width=47)
  Sort Key: (sum(o.actual_price)) DESC
  -> HashAggregate (cost=3578.50..3578.81 rows=31 width=47)
    Group By Key: ad.province
    -> Hash Join (cost=3340.21..3571.74 rows=1351 width=15)
      Hash Cond: (ad.address_key = o.address_key)
      -> Seq Scan on address_dimension ad (cost=0.00..188.02 rows=8002 width=14)
      -> Hash (cost=3323.32..3323.32 rows=1351 width=9)
        -> Hash Join (cost=17.50..3323.32 rows=1351 width=9)
          Hash Cond: (o.add_date = dd.date_key)
          -> Seq Scan on litemall_orders o (cost=0.00..3041.00 rows=100000 width=13)
          -> Hash (cost=17.53..17.53 rows=2 width=4)
            -> Index Scan using <16508>btree_date_dimension_year on date_dimension dd (cost=0.00..17.53 rows=2 width=4)
              Index Cond: (year = 2020)
              Filter: ((month)::bigint = 3)

(15 rows)

tpch=#

```

步骤 11 【附加题】有兴趣的同学可以尝试并截图记录于此。仅需要从 queries.sql 文件里选择一条或多条进行索引优化即可。

(截图执行语句和结果)

选择 queries.sql 中的如下语句：

```
tpch=# select
    l_returnflag,
    l_linestatus,
    sum(l_quantity) as sum_qty,
    sum(l_extendedprice) as sum_base_price,
    sum(l_extendedprice * (1 - l_discount)) as sum_disc_price,
    sum(l_extendedprice * (1 - l_discount) * (1 + l_tax)) as sum_charge,
    avg(l_quantity) as avg_qty,
    avg(l_extendedprice) as avg_price,
    avg(l_discount) as avg_disc,
    count(*) as count_order
from
    lineitem
where
    l_shipdate <= date '1998-12-01' - interval '90' day
group by
    l_returnflag,
    l_linestatus
order by
    l_returnflag,
    l_linestatus;
```

耗时较长，执行结果如下：

```
tpch=# select
tpch=# l_returnflag,
tpch=# l_linestatus,
tpch=# sum(l_quantity) as sum_qty,
tpch=# sum(l_extendedprice) as sum_base_price,
tpch=# sum(l_extendedprice * (1 - l_discount)) as sum_disc_price,
tpch=# sum(l_extendedprice * (1 - l_discount) * (1 + l_tax)) as sum_charge,
tpch=# avg(l_quantity) as avg_qty,
tpch=# avg(l_extendedprice) as avg_price,
tpch=# avg(l_discount) as avg_disc,
tpch=# count(*) as count_order
tpch=# from
tpch=# lineitem
tpch=# where
tpch=# l_shipdate <= date '1998-12-01' - interval '90' day
tpch=# group by
tpch=# l_returnflag,
tpch=# l_linestatus
tpch=# order by
tpch=# l_returnflag,
tpch=# l_linestatus;

 l_returnflag | l_linestatus | sum_qty | sum_base_price | sum_disc_price | sum_charge | avg_qty | avg_price |
-----+-----+-----+-----+-----+-----+-----+-----+
A | F | 37734107.00 | 56586554400.73 | 53758257134.8700 | 55909065222.827692 | 25.5220058532573370 | 38275.129734621672 |
04998529583839761162 | 1478493
N | F | 991417.00 | 1487504710.38 | 1413082168.0541 | 1469649223.194375 | 25.5164719205229835 | 38284.467760848304 |
05009342667421629691 | 38854
N | O | 74476040.00 | 111701729697.74 | 106118236307.6056 | 110567043072.497010 | 25.5022267695849915 | 38249.117988908270 |
04999658065370408037 | 2920374
R | F | 37719753.00 | 56560041300.90 | 53741202684.6040 | 55889619119.831932 | 25.5057936126907707 | 38250.854626099657 |
05000940583012705647 | 1478870
(4 rows)

tpch=#
```

使用 explain，对该 SQL 加以分析

```
tpch=# EXPLAIN
select
    l_returnflag,
    l_linestatus,
```

```

sum(l_quantity) as sum_qty,
sum(l_extendedprice) as sum_base_price,
sum(l_extendedprice * (1 - l_discount)) as sum_disc_price,
sum(l_extendedprice * (1 - l_discount) * (1 + l_tax)) as sum_charge,
avg(l_quantity) as avg_qty,
avg(l_extendedprice) as avg_price,
avg(l_discount) as avg_disc,
count(*) as count_order
from
    lineitem
where
    l_shipdate <= date '1998-12-01' - interval '90' day
group by
    l_returnflag,
    l_linestatus
order by
    l_returnflag,
    l_linestatus;

```

获得执行计划结果为：

```

tpch=#
tpch=# EXPLAIN
tpch=# select
tpch=#   l_returnflag,
tpch=#   l_linestatus,
tpch=#   sum(l_quantity) as sum_qty,
tpch=#   sum(l_extendedprice) as sum_base_price,
tpch=#   sum(l_extendedprice * (1 - l_discount)) as sum_disc_price,
tpch=#   sum(l_extendedprice * (1 - l_discount) * (1 + l_tax)) as sum_charge,
tpch=#   avg(l_quantity) as avg_qty,
tpch=#   avg(l_extendedprice) as avg_price,
tpch=#   avg(l_discount) as avg_disc,
tpch=#   count(*) as count_order
tpch=# from
tpch=#   lineitem
tpch=# where
tpch=#   l_shipdate <= date '1998-12-01' - interval '90' day
tpch=# group by
tpch=#   l_returnflag,
tpch=#   l_linestatus
tpch=# order by
tpch=#   l_returnflag,
tpch=#   l_linestatus;

----- QUERY PLAN -----
Sort  (cost=436259.85..436259.86 rows=6 width=257)
  Sort Key: l_returnflag, l_linestatus
  -> HashAggregate  (cost=436259.66..436259.77 rows=6 width=257)
    Group By Key: l_returnflag, l_linestatus
    -> Seq Scan on lineitem  (cost=0.00..199621.06 rows=5915965 width=25)
        Filter: (l_shipdate <= '1998-09-02 00:00:00'::timestamp without time zone)
(6 rows)

tpch=#

```

使用索引推荐功能，对查询语句进行推荐：

```

tpch=# select * from gs_index_advise('
select
    l_returnflag,
    l_linestatus,
    sum(l_quantity) as sum_qty,

```



```

sum(l_extendedprice) as sum_base_price,
sum(l_extendedprice * (1 - l_discount)) as sum_disc_price,
sum(l_extendedprice * (1 - l_discount) * (1 + l_tax)) as sum_charge,
avg(l_quantity) as avg_qty,
avg(l_extendedprice) as avg_price,
avg(l_discount) as avg_disc,
count(*) as count_order
from
    lineitem
where
    l_shipdate <= date "1998-12-01" - interval "90" day
group by
    l_returnflag,
    l_linestatus
order by
    l_returnflag,
    l_linestatus');

```

获得索引推荐结果如下：

```

tpch=# select * from gs_index_advise('
tpch'# select
tpch'# l_returnflag,
tpch'# l_linestatus,
tpch'# sum(l_quantity) as sum_qty,
tpch'# sum(l_extendedprice) as sum_base_price,
tpch'# sum(l_extendedprice * (1 - l_discount)) as sum_disc_price,
tpch'# sum(l_extendedprice * (1 - l_discount) * (1 + l_tax)) as sum_charge,
tpch'# avg(l_quantity) as avg_qty,
tpch'# avg(l_extendedprice) as avg_price,
tpch'# avg(l_discount) as avg_disc,
tpch'# count(*) as count_order
tpch'# from
tpch'# lineitem
tpch'# where
tpch'# l_shipdate <= date ''1998-12-01'' - interval ''90'' day
tpch'# group by
tpch'# l_returnflag,
tpch'# l_linestatus
tpch'# order by
tpch'# l_returnflag,
tpch'# l_linestatus');
 schema | table | column
-----+-----+-----
 public | lineitem | (l_returnflag,l_linestatus)
(1 row)

```

在 lineitem 表上创建虚拟索引列。

```

tpch=# select * from hypopg_create_index('create index on
lineitem(l_returnflag,l_linestatus)');

```

```

tpch=# select * from hypopg_create_index('create index on lineitem(l_returnflag,l_linestatus)');
 indexrelid | indexname
-----+-----
 16514 | <16514>btree_lineitem_l_returnflag_l_linestatus
(1 row)

```

查看创建的虚拟索引列。

```
tpch=# select * from hypopg_display_index();
```

```
tpch=# select * from hypopg_display_index();
      indexname      | indexrelid | table  | column
-----+-----+-----+-----
<16514>btree_lineitem_l_returnflag_l_linestatus |      16514 | lineitem | (l_returnflag, l_linestatus)
(1 row)
```

获取索引虚拟列大小结果

```
tpch=# select * from hypopg_estimate_size(16514);
```

返回结果为：

```
tpch=# select * from hypopg_estimate_size(16514);
hypopg_estimate_size
-----
                219709440
(1 row)
```

开启 GUC 参数 enable_hypo_index。

```
tpch=# set enable_hypo_index = on;
```

再次使用 explain，对刚才执行的 SQL 加以分析

```
tpch=# set enable_hypo_index = on;
SET
tpch=# EXPLAIN
tpch=# select
tpch=# l_returnflag,
tpch=# l_linestatus,
tpch=# sum(l_quantity) as sum_qty,
tpch=# sum(l_extendedprice) as sum_base_price,
tpch=# sum(l_extendedprice * (1 - l_discount)) as sum_disc_price,
tpch=# sum(l_extendedprice * (1 - l_discount) * (1 + l_tax)) as sum_charge,
tpch=# avg(l_quantity) as avg_qty,
tpch=# avg(l_extendedprice) as avg_price,
tpch=# avg(l_discount) as avg_disc,
tpch=# count(*) as count_order
tpch=# from
tpch=# lineitem
tpch=# where
tpch=# l_shipdate <= date '1998-12-01' - interval '90' day
tpch=# group by
tpch=# l_returnflag,
tpch=# l_linestatus
tpch=# order by
tpch=# l_returnflag,
tpch=# l_linestatus;

QUERY PLAN
-----
Sort (cost=436259.85..436259.86 rows=6 width=257)
  Sort Key: l_returnflag, l_linestatus
    -> HashAggregate (cost=436259.66..436259.77 rows=6 width=257)
      Group By Key: l_returnflag, l_linestatus
      -> Seq Scan on lineitem (cost=0.00..199621.06 rows=5915965 width=25)
        Filter: (l_shipdate <= '1998-09-02 00:00:00'::timestamp without time zone)
(6 rows)

tpch=#
```

可以看出该语句并未奏效，分析发现，这是由于该语句仅是单表操作。

清理创建的索引虚拟列：

```
tpch=# select * from hypopg_reset_index();
```

```
tpch=# select * from hypopg_reset_index();
hypopg_reset_index
-----
(1 row)
```

因此挑选多表操作的语句：

```
tpch=# select
      s_acctbal,
      s_name,
      n_name,
      p_partkey,
      p_mfgr,
      s_address,
      s_phone,
      s_comment
from
      part,
      supplier,
      partsupp,
      nation,
      region
where
      p_partkey = ps_partkey
      and s_suppkey = ps_suppkey
      and p_size = 15
      and p_type like '%BRASS'
      and s_nationkey = n_nationkey
      and n_regionkey = r_regionkey
      and r_name = 'EUROPE'
      and ps_supplycost = (
          select
              min(ps_supplycost)
          from
              partsupp,
              supplier,
              nation,
              region
          where
p_partkey = ps_partkey
              and s_suppkey = ps_suppkey
              and s_nationkey = n_nationkey
              and n_regionkey = r_regionkey
              and r_name = 'EUROPE'
          )
order by
      s_acctbal desc,
      n_name,
      s_name,
      p_partkey;
;
```

```

tpch=# and r_name = 'EUROPE'
tpch=# and ps_supplycost = (
tpch=# select
tpch=# min(ps_supplycost)
tpch=# from
tpch=# partsupp,
tpch=# supplier,
tpch=# nation,
tpch=# region
tpch=# where
tpch=# p_partkey = ps_partkey
tpch=# and s_suppkey = ps_suppkey
tpch=# and s_nationkey = n_nationkey
tpch=# and n_regionkey = r_regionkey
tpch=# and r_name = 'EUROPE'
tpch=# )
tpch=# order by
tpch=# s_acctbal desc,
tpch=# n_name,
tpch=# s_name,
tpch=# p_partkey;
tpch=#
s_acctbal | s_name | n_name | p_partkey | p_mfgr | s_address
| s_phone | s_comment
-----
9038.53 | Supplier#000005359 | UNITED KINGDOM | 185358 | Manufacturer#4 | QKwHYh,vZG1wu2FWEJqLDx04
| 33-429-790-6131 | uriously regular requests hag
9037.84 | Supplier#000005069 | ROMANIA | 188438 | Manufacturer#1 | ANDENSOSmk,m1q23Xfb5RwT6dvUcvt6Qs
| 28-598-602-3537 | efully express instructions. regular requests against the slyly fin
9036.22 | Supplier#000005259 | UNITED KINGDOM | 249 | Manufacturer#4 | B3rcpBxbSE1wMey2RH J
| 33-328-228-2057 | etect about the furiously final accounts. slyly ironic pinto beans sleep inside the furiously
9023.77 | Supplier#000002324 | GERMANY | 29821 | Manufacturer#4 | y30D9Uyw5T0k
| 17-770-200-1839 | ackages boost blithely, blithely regular deposits c
9071.22 | Supplier#000006373 | GERMANY | 43868 | Manufacturer#5 | J8fcXwTqM
| 17-813-486-8637 | etect blithely bold asymptotes. fluffily ironic platelets wake furiously; blit
9878.78 | Supplier#000001286 | GERMANY | 81285 | Manufacturer#2 | YKA,E2fj;Vd7eUrzp2Ef8j1QxGo2DFnosaTEH
| 17-516-924-4574 | regular accounts. furiously unusual courts above the fi
9878.78 | Supplier#000001286 | GERMANY | 81285 | Manufacturer#2 | YKA,E2fj;Vd7eUrzp2Ef8j1QxGo2DFnosaTEH
| 17-516-924-4574 | regular accounts. furiously unusual courts above the fi
9852.52 | Supplier#000009073 | RUSSIA | 18972 | Manufacturer#2 | t5L67YdBYTH6o,Vz24jpDyQ0
| 32-188-594-7838 | rns wake final foxes. carefully unusual depende
9847.85 | Supplier#000008897 | RUSSIA | 189587 | Manufacturer#2 | xM97bpE6QNZdwLoX
| 32-375-640-3593 | the special excuses. silent sentiments serve carefully final ac
9847.57 | Supplier#000006345 | FRANCE | 88344 | Manufacturer#1 | VSt3rzK3qG698u6Ld8HhO8yvrTcSTsvQLDQ0ag
| 16-886-768-7945 | ges. slyly regular requests are. ruthless, express excuses cajole blithely across the unu
9847.57 | Supplier#000006345 | FRANCE | 173827 | Manufacturer#2 | VSt3rzK3qG698u6Ld8HhO8yvrTcSTsvQLDQ0ag
| 16-886-768-7945 | ges. slyly regular requests are. ruthless, express excuses cajole blithely across the unu
9836.93 | Supplier#000007342 | RUSSIA | 4841 | Manufacturer#4 | J0LX7C1,7xrtZSS0w
| 32-399-414-5388 | blithely carefully bold theodolites. fur
9817.18 | Supplier#000002352 | RUSSIA | 124815 | Manufacturer#2 | 4Lf0HUZjg;EbAKw TgdKcg0c4D4uCYw
| 32-551-831-1437 | wake carefully alongside of the carefully final ex
9817.18 | Supplier#000002352 | RUSSIA | 152351 | Manufacturer#3 | 4Lf0HUZjg;EbAKw TgdKcg0c4D4uCYw
| 32-551-831-1437 | wake carefully alongside of the carefully final ex
9739.86 | Supplier#000003384 | FRANCE | 138357 | Manufacturer#2 | o,Z3v4P0sf0vE k9U1b 63lucX,I
| 16-404-013-5025 | s after the furiously bold packages sleep fluffily idly final requests: quickly final
9721.95 | Supplier#000008757 | UNITED KINGDOM | 158241 | Manufacturer#3 | Atg6nM4dT2

```

使用 explain，对该 SQL 加以分析

```

tpch=# EXPLAIN
select
s_acctbal,
s_name,
n_name,
p_partkey,
p_mfgr,
s_address,
s_phone,
s_comment
from
part,
supplier,
partsupp,
nation,
region
where
p_partkey = ps_partkey
and s_suppkey = ps_suppkey
and p_size = 15
and p_type like '%BRASS'
and s_nationkey = n_nationkey
and n_regionkey = r_regionkey
and r_name = 'EUROPE'
and ps_supplycost = (
select
min(ps_supplycost)
from
partsupp,

```

```

supplier,
nation,
region

where

p_partkey = ps_partkey

and s_suppkey = ps_suppkey
and s_nationkey = n_nationkey
and n_regionkey = r_regionkey
and r_name = 'EUROPE'

)
order by
s_acctbal desc,
n_name,
s_name,
p_partkey;
;
```

获得执行计划结果为：

```

tpch=# s_acctbal desc,
tpch=# n_name,
tpch=# s_name,
tpch=# p_partkey;

QUERY PLAN

Sort (cost=71159.24..71159.24 rows=1 width=270)
Sort Key: public.supplier.s_acctbal DESC, public.nation.n_name, public.supplier.s_name, public.part.p_partkey
-> Hash Join (cost=42631.34..71159.23 rows=1 width=270)
Hash Cond: ((public.part.p_partkey = subquery."?column?") AND (public.partsupp.ps_supplycost = subquery.min))
-> Hash Join (cost=7490.82..36027.56 rows=17 width=280)
Hash Cond: (public.partsupp.ps_suppkey = public.supplier.s_suppkey)
-> Hash Join (cost=7115.00..35631.60 rows=2951 width=44)
Hash Cond: (public.partsupp.ps_partkey = public.part.p_partkey)
-> Seq Scan on partsupp (cost=0.00..25487.00 rows=800000 width=14)
-> Hash (cost=7106.00..7106.00 rows=727 width=30)
Filter: ((p_type)::text ~~ '%BRASS'::text) AND (p_size = 15))
-> Hash (cost=384.03..384.03 rows=56 width=244)
-> Hash Join (cost=24.46..384.03 rows=56 width=244)
Hash Cond: (public.supplier.s_nationkey = public.nation.n_nationkey)
-> Seq Scan on supplier (cost=0.00..323.00 rows=10000 width=144)
-> Hash (cost=24.45..24.45 rows=1 width=100)
-> Hash Join (cost=12.24..24.45 rows=1 width=100)
Hash Cond: (public.nation.n_regionkey = public.region.r_regionkey)
-> Seq Scan on nation (cost=0.00..11.76 rows=176 width=112)
-> Hash (cost=12.22..12.22 rows=1 width=4)
-> Seq Scan on region (cost=0.00..12.22 rows=1 width=4)
Filter: (r_name = 'EUROPE'::bpchar)
-> Hash (cost=35131.28..35131.28 rows=16 width=36)
-> Subquery Scan on subquery (cost=35130.06..35131.28 rows=16 width=36)
-> HashAggregate (cost=35130.06..35131.12 rows=16 width=42)
Group By Key: public.partsupp.ps_partkey
-> Hash Join (cost=7490.82..35130.88 rows=16 width=10)
Hash Cond: (public.partsupp.ps_suppkey = public.supplier.s_suppkey)
-> Hash Semi Join (cost=7115.00..34734.02 rows=2951 width=14)
Hash Cond: (public.partsupp.ps_partkey = public.part.p_partkey)
-> Seq Scan on partsupp (cost=0.00..25487.00 rows=800000 width=14)
-> Hash (cost=7106.00..7106.00 rows=727 width=4)
-> Seq Scan on part (cost=0.00..7106.00 rows=727 width=4)
Filter: ((p_type)::text ~~ '%BRASS'::text) AND (p_size = 15))
-> Hash (cost=384.03..384.03 rows=56 width=4)
-> Hash Join (cost=24.46..384.03 rows=56 width=4)
Hash Cond: (public.supplier.s_nationkey = public.nation.n_nationkey)
-> Seq Scan on supplier (cost=0.00..323.00 rows=10000 width=8)
-> Hash (cost=24.45..24.45 rows=1 width=4)
-> Hash Join (cost=12.24..24.45 rows=1 width=4)
Hash Cond: (public.nation.n_regionkey = public.region.r_regionkey)
-> Seq Scan on nation (cost=0.00..11.76 rows=176 width=8)
-> Hash (cost=12.22..12.22 rows=1 width=4)
-> Seq Scan on region (cost=0.00..12.22 rows=1 width=4)
Filter: (r_name = 'EUROPE'::bpchar)

(46 rows)

tpch=#
```

使用索引推荐功能，对查询语句进行推荐：

```
tpch=# select * from gs_index_advise(' select
```



```

        s_acctbal,
        s_name,
        n_name,
        p_partkey,
        p_mfgr,
        s_address,
        s_phone,
        s_comment
from
    part,
    supplier,
    partsupp,
    nation,
    region
where
    p_partkey = ps_partkey
    and s_suppkey = ps_suppkey
    and p_size = 15
    and p_type like "%BRASS"
    and s_nationkey = n_nationkey
    and n_regionkey = r_regionkey
    and r_name = "EUROPE"
    and ps_supplycost = (
        select
            min(ps_supplycost)
        from
            partsupp,
            supplier,
            nation,
            region
        where
p_partkey = ps_partkey
            and s_suppkey = ps_suppkey
            and s_nationkey = n_nationkey
            and n_regionkey = r_regionkey
            and r_name = "EUROPE"
        )
order by
    s_acctbal desc,
    n_name,
    s_name,
    p_partkey;
');

```

获得索引推荐结果如下：

```

tpch'#      region
tpch'# where
tpch'#      p_partkey = ps_partkey
tpch'#      and s_suppkey = ps_suppkey
tpch'#      and p_size = 15
tpch'#      and p_type like '%BRASS'
tpch'#      and s_nationkey = n_nationkey
tpch'#      and n_regionkey = r_regionkey
tpch'#      and r_name = 'EUROPE'
tpch'#      and ps_supplycost = (
tpch'#          select
tpch'#              min(ps_supplycost)
tpch'#          from
tpch'#              partsupp,
tpch'#              supplier,
tpch'#              nation,
tpch'#              region
tpch'#          where
tpch'#      p_partkey = ps_partkey
tpch'#          and s_suppkey = ps_suppkey
tpch'#          and s_nationkey = n_nationkey
tpch'#          and n_regionkey = r_regionkey
tpch'#          and r_name = 'EUROPE'
tpch'#      )
tpch'# order by
tpch'#      s_acctbal desc,
tpch'#      n_name,
tpch'#      s_name,
tpch'#      p_partkey;
tpch'# ');
  schema | table  | column
-----+-----+-----
public  | part   | (p_partkey,p_size)
public  | supplier | (s_suppkey,s_nationkey),(s_acctbal,s_name)
public  | partsupp | (ps_partkey,ps_suppkey)
public  | nation | (n_name)
public  | region | (r_regionkey)
(5 rows)

tpch=# 

```

在 part、supplier、partsupp、nation 和 region 表上创建虚拟索引列：

```

tpch=# select * from hypopg_create_index('create index on part(p_partkey,p_size)');
tpch=# select * from hypopg_create_index('create index on supplier(s_suppkey,s_nationkey,
s_acctbal,s_name)');
tpch=# select * from hypopg_create_index('create index on
partsupp(ps_partkey,ps_suppkey)');
tpch=# select * from hypopg_create_index('create index on nation(n_name)');
tpch=# select * from hypopg_create_index('create index on region(r_regionkey)');

```

```

tpch=# select * from hypopg_create_index('create index on part(p_partkey,p_size)');
indexrelid |          indexname
-----+-----
16515 | <16515>btree_part_p_partkey_p_size
(1 row)

tpch=# select * from hypopg_create_index('create index on supplier(s_suppkey,s_nationkey, s_acctbal,s_name)');
indexrelid |          indexname
-----+-----
16516 | <16516>btree_supplier_s_suppkey_s_nationkey_s_acctbal_s_name
(1 row)

tpch=# select * from hypopg_create_index('create index on partsupp(ps_partkey,ps_suppkey)');
indexrelid |          indexname
-----+-----
16517 | <16517>btree_partsupp_ps_partkey_ps_suppkey
(1 row)

tpch=# select * from hypopg_create_index('create index on nation(n_name)');
indexrelid |          indexname
-----+-----
16518 | <16518>btree_nation_n_name
(1 row)

tpch=# select * from hypopg_create_index('create index on region(r_regionkey)');
indexrelid |          indexname
-----+-----
16519 | <16519>btree_region_r_regionkey
(1 row)

```

查看创建的虚拟索引列：

```
tpch=# select * from hypopg_display_index();
```

获得创建虚拟列的结果：

```

tpch=# select * from hypopg_display_index();
indexname | indexrelid | table | column
-----+-----+-----+-----
<16515>btree_part_p_partkey_p_size | 16515 | part | (p_partkey, p_size)
<16516>btree_supplier_s_suppkey_s_nationkey_s_acctbal_s_name | 16516 | supplier | (s_suppkey, s_nationkey, s_acctbal, s_name)
<16517>btree_partsupp_ps_partkey_ps_suppkey | 16517 | partsupp | (ps_partkey, ps_suppkey)
<16518>btree_nation_n_name | 16518 | nation | (n_name)
<16519>btree_region_r_regionkey | 16519 | region | (r_regionkey)
(5 rows)

```

获取索引虚拟列大小结果

```

tpch=# select * from hypopg_estimate_size(16515);
tpch=# select * from hypopg_estimate_size(16516);
tpch=# select * from hypopg_estimate_size(16517);
tpch=# select * from hypopg_estimate_size(16518);
tpch=# select * from hypopg_estimate_size(16519);

```

返回结果为：

```

tpch=# select * from hypopg_estimate_size(16515);
hypopg_estimate_size
-----
8364032
(1 row)

tpch=# select * from hypopg_estimate_size(16516);
hypopg_estimate_size
-----
1146880
(1 row)

tpch=# select * from hypopg_estimate_size(16517);
hypopg_estimate_size
-----
33472512
(1 row)

tpch=# select * from hypopg_estimate_size(16518);
hypopg_estimate_size
-----
8192
(1 row)

tpch=# select * from hypopg_estimate_size(16519);
hypopg_estimate_size
-----
8192
(1 row)

```

开启 GUC 参数 enable_hypo_index：

```
tpch=# set enable_hypo_index = on;
```

再次使用 explain , 对刚才执行的 SQL 加以分析

```
tpch=# EXPLAIN
select
    s_acctbal,
    s_name,
    n_name,
    p_partkey,
    p_mfgr,
    s_address,
    s_phone,
    s_comment
from
    part,
    supplier,
    partsupp,
    nation,
    region
where
    p_partkey = ps_partkey
    and s_suppkey = ps_suppkey
    and p_size = 15
    and p_type like '%BRASS'
    and s_nationkey = n_nationkey
    and n_regionkey = r_regionkey
    and r_name = 'EUROPE'
    and ps_supplycost = (
        select
            min(ps_supplycost)
        from
            partsupp,
            supplier,
            nation,
            region
        where
            p_partkey = ps_partkey
            and s_suppkey = ps_suppkey
            and s_nationkey = n_nationkey
            and n_regionkey = r_regionkey
            and r_name = 'EUROPE'
        )
order by
    s_acctbal desc,
    n_name,
    s_name,
    p_partkey;
;
```

```

tpch=# p_partkey;
QUERY PLAN
-----
Sort (cost=6461.31..6461.31 rows=1 width=270)
Sort Key: public.supplier.s_acctbal DESC, public.nation.n_name, public.supplier.s_name, public.part.p_partkey
-> Nested Loop (cost=6485.86..6461.30 rows=1 width=270)
-> Hash Join (cost=6485.86..6460.95 rows=1 width=274)
Hash Cond: (public.supplier.s_nationkey = public.nation.n_nationkey)
-> Nested Loop (cost=6391.10..6446.98 rows=1 width=178)
-> Nested Loop (cost=6391.10..6446.78 rows=1 width=34)
Join Filter: ((min(public.partsupp.ps_supplycost)) = public.partsupp.ps_supplycost)
-> Nested Loop (cost=6391.10..6436.30 rows=3 width=66)
-> HashAggregate (cost=6391.10..6391.26 rows=16 width=42)
Group By Key: public.partsupp.ps_partkey
-> Hash Join (cost=3859.74..6391.02 rows=16 width=10)
Hash Cond: (public.partsupp.ps_suppkey = public.supplier.s_suppkey)
-> Nested Loop (cost=3748.28..6268.32 rows=2951 width=14)
-> HashAggregate (cost=3748.28..3755.55 rows=727 width=4)
Group By Key: public.part.p_partkey
-> Index Scan using <16515>btree_part_p_partkey_p_size on part (cost=0.00..3746.46 rows=72
7 width=4)
Index Cond: (p_size = 15)
Filter: ((p_type)::text ~~ 'NBRASS'::text)
-> Index Scan using <16517>btree_partsupp_ps_partkey_ps_suppkey on partsupp (cost=0.00..3.42 row
s=4 width=14)
Index Cond: (ps_partkey = public.part.p_partkey)
-> Hash (cost=110.76..110.76 rows=56 width=4)
-> Nested Loop (cost=12.24..110.76 rows=56 width=4)
-> Hash Join (cost=12.24..24.45 rows=1 width=4)
Hash Cond: (public.nation.n_regionkey = public.region.r_regionkey)
-> Seq Scan on nation (cost=0.00..11.76 rows=176 width=8)
-> Hash (cost=12.22..12.22 rows=1 width=4)
-> Seq Scan on region (cost=0.00..12.22 rows=1 width=4)
Filter: (r_name = 'EUROPE'::bpchar)
-> Index Only Scan using <16516>btree_supplier_s_suppkey_s_nationkey_s_acctbal_s_name on su
pplier (cost=0.00..82.31 rows=400 width=0)
Index Cond: (s_nationkey = public.nation.n_nationkey)
-> Index Scan using <16515>btree_part_p_partkey_p_size on part (cost=0.00..2.27 rows=53 width=30)
Index Cond: ((p_partkey = public.partsupp.ps_partkey) AND (p_size = 15))
Filter: ((p_type)::text ~~ 'NBRASS'::text)
-> Index Scan using <16517>btree_partsupp_ps_partkey_ps_suppkey on partsupp (cost=0.00..3.42 rows=4 width=14)
Index Cond: (ps_partkey = public.part.p_partkey)
-> Index Scan using <16516>btree_supplier_s_suppkey_s_nationkey_s_acctbal_s_name on supplier (cost=0.00..0.27 rows=1 width=144)
)
Index Cond: (s_suppkey = public.partsupp.ps_suppkey)
-> Hash (cost=11.76..11.76 rows=176 width=112)
-> Seq Scan on nation (cost=0.00..11.76 rows=176 width=112)
-> Index Scan using <16510>btree_region_r_regionkey on region (cost=0.00..0.34 rows=1 width=4)
Index Cond: (r_regionkey = public.nation.n_regionkey)
Filter: (r_name = 'EUROPE'::bpchar)
(43 rows)

```

结果中出现：

```

-> Index Scan using <16517>btree_partsupp_ps_partkey_ps_suppkey on partsupp (cost=0.00..3.42 rows=4 width=14)
Index Cond: (ps_partkey = public.part.p_partkey)

-> Index Scan using <16510>btree_region_r_regionkey on region (cost=0.00..0.34 rows=1 width=4)
Index Cond: (r_regionkey = public.nation.n_regionkey)

```

说明索引推荐已经奏效。

清理创建的索引虚拟列：

```

tpch=# select * from hypopg_reset_index();

tpch=# select * from hypopg_reset_index();
hypopg_reset_index
-----
(1 row)

```

关卡四【附加题】： openGauss 的 DB4AI 特性应用

*本关卡为附加题，有兴趣的同学可以尝试实验并记录于此。

1. 关卡验证

步骤 10 利用训练好的逻辑回归模型预测数据，并与 SVM 算法进行比较，将执行结果截图。


```
openGauss=# SELECT tax, bath, size, price, price < 100000 AS price_actual, PREDICT BY
house_binary_classifier (FEATURES tax, bath, size) AS price_svm_pred, PREDICT BY
house_logistic_classifier (FEATURES tax, bath, size) AS price_logistic_pred FROM houses;
```

（截图执行语句和结果）

```
openGauss=# SELECT tax, bath, size, price, price < 100000 AS price_actual, PREDICT BY house_binary_classif
_pred, PREDICT BY house_logistic_classifier (FEATURES tax, bath, size) AS price_logistic_pred FROM houses;
tax | bath | size | price | price_actual | price_svm_pred | price_logistic_pred
-----+-----+-----+-----+-----+-----+-----
590 | 1 | 770 | 50000 | t | t | t
1050 | 2 | 1410 | 85000 | t | t | t
20 | 1 | 1060 | 22500 | t | t | t
870 | 2 | 1300 | 90000 | t | t | t
1320 | 2 | 1500 | 133000 | f | t | t
1350 | 1 | 820 | 90500 | t | f | f
2790 | 2.5 | 2130 | 260000 | f | f | f
680 | 1 | 1170 | 142500 | f | t | t
1840 | 2 | 1500 | 160000 | f | f | f
3680 | 2 | 2790 | 240000 | f | f | f
1660 | 1 | 1030 | 87000 | t | f | f
1620 | 2 | 1250 | 118600 | f | f | f
3100 | 2 | 1760 | 140000 | f | f | f
2070 | 3 | 1550 | 148000 | f | f | f
650 | 1.5 | 1450 | 65000 | t | t | t
(15 rows)
```

可以看出，逻辑回归模型与 SVM 算法的结果类似，说明某些样例较为困难，这两种方法都无法预测正确。

清理工作：资源释放

1. 关卡验证

步骤 3 查看到列表中已没有资源时，表示弹性云服务器已删除。

（截图执行语句和结果）

